

WEB3 WEBSITE COPY SAMPLE

Decentralized AI Infrastructure

Production-style website copy for a fictional decentralized AI inference network

Prepared by Firman Siregar

Web3 Security Researcher | DeFi, Smart Contracts & Risk Analysis

Purpose

A portfolio sample demonstrating research-backed Web3 website copy: clear technical communication, credible claims, practical explanations, and non-shill positioning for a decentralized infrastructure product.

Note: This is a fictional portfolio project. Product names, metrics, economics, and technical claims are illustrative and should not be interpreted as claims about a real protocol.

01 — HOME

Website copy designed to introduce a decentralized AI inference network to a technically curious but non-specialist audience.

HERO HEADLINE

Run AI where the network runs it.

SUBHEADLINE

A decentralized inference network that connects developers with distributed compute—without putting every workload behind a single infrastructure provider.

PRIMARY CTA

Explore the network

SECONDARY CTA

Build with the API

SECTION — THE NEXT INFRASTRUCTURE LAYER

The internet decentralized information. BitTorrent decentralized distribution. AI is creating a new demand for compute.

The next question is not only who builds AI models, but who provides the infrastructure that runs them. A decentralized inference network coordinates independent compute providers so applications can request AI inference from a distributed pool of resources.

SECTION — WHY DECENTRALIZED INFERENCE?

AI inference can be infrastructure-intensive, geographically distributed, and sensitive to availability and cost. A decentralized model changes the infrastructure relationship: instead of relying on one centralized provider, applications can interact with a network of independent participants.

The goal is not decentralization for its own sake. It is to create an infrastructure layer where supply can come from multiple operators, developers can consume compute through a common interface, and network rules can coordinate participation.

SECTION — FROM BITCOIN TO BITTORRENT TO DECENTRALIZED AI

Bitcoin demonstrated that value transfer could operate without a central clearing institution. BitTorrent demonstrated that large-scale distribution could be coordinated across independent participants. Decentralized AI infrastructure applies a related idea to compute: coordinate a distributed supply of resources around a common protocol.

SECTION — FOR DEVELOPERS

Integrate AI inference without rebuilding your infrastructure around a single compute provider. Use a standardized API to submit workloads, receive responses, and build applications on top of the network.

CTA

Start building

SECTION — FOR HOSTS

Put available compute to work. Hosts contribute infrastructure to the network and receive compensation according to the network's rules, work performed, and applicable market conditions.

CTA

Become a host

02 — HOSTS

Practical copy for infrastructure providers considering participation in the network.

HERO HEADLINE

Turn available compute into network capacity.

SUBHEADLINE

Contribute compatible hardware, serve inference workloads, and participate in an open compute marketplace.

WHAT HOSTS DO

Hosts provide compute resources that can be used to process inference requests. The network coordinates workload assignment, request handling, and settlement according to its protocol rules.

ECONOMICS

Host economics depend on factors such as available hardware, workload demand, utilization, operating costs, and network conditions. Actual returns are not guaranteed. Before participating, operators should evaluate their own hardware, energy, bandwidth, maintenance, and opportunity costs.

SETUP

1. Prepare compatible hardware and networking.
2. Install the required host software.
3. Configure the node according to network requirements.
4. Connect to the network and begin accepting eligible workloads.
5. Monitor utilization, performance, and operating costs.

WHY PARTICIPATE?

A distributed compute network creates another way for infrastructure operators to put hardware capacity to productive use. For hosts, the key questions are practical: What hardware is required? What workloads are available? How is work measured? How are rewards calculated? What does operating the infrastructure actually cost?

CTA

Check host requirements

03 — DEVELOPERS

Copy focused on integration, API flow, and practical application.

HERO HEADLINE

Build AI applications on distributed compute.

SUBHEADLINE

A common interface for sending inference workloads to a decentralized network of compute providers.

HOW IT WORKS

Your application sends an inference request through the network interface. The network routes the request to an eligible provider, the provider processes the workload, and the result is returned to your application.

The exact routing, verification, settlement, and availability mechanisms depend on the network implementation and its protocol rules.

API FLOW

- 01 — Create your application credentials
- 02 — Submit an inference request
- 03 — Network coordinates an eligible provider
- 04 — Provider processes the workload
- 05 — Receive the inference result
- 06 — Monitor usage and application performance

USE CASES

Decentralized AI infrastructure can support applications such as AI assistants, autonomous agents, developer tools, content workflows, model-powered applications, and other products that require programmatic inference.

BUILD WITHOUT THE HYPE

Infrastructure should be evaluated by what it actually does. Start with the API, understand the execution model, measure performance, and evaluate cost and reliability against your application's requirements.

CTA

Read the developer documentation

04 — OPTIONAL TOKEN MODULE

Neutral informational copy deliberately separated from the core product narrative.

HEADING

Network token

BODY

The network may use a native token as part of its protocol design. Depending on the implementation, the token can be used for functions such as network participation, settlement, governance, or other protocol-level mechanisms.

The specific role, supply, distribution, utility, and governance rights should be defined by the project's official documentation. Nothing in this website copy should be interpreted as an offer, recommendation, or representation of investment value.

CTA

Read the token documentation

05 — WRITING APPROACH

The sample intentionally demonstrates the writing principles relevant to technical Web3 projects.

Technical without unnecessary jargon	Technical concepts are introduced in plain language and expanded only where additional detail is useful.
No invented metrics	No fictional TPS, revenue, APY, latency, market-share, or ROI claims are presented as product facts.
No investment framing	The optional token section uses informational language rather than price, return, or investment claims.

Audience separation	Home establishes the narrative; Hosts focuses on operations and economics; Developers focuses on integration and use cases.
Research-first	The copy is structured so product-specific claims can be replaced with verified facts from documentation before publication.

Portfolio positioning: This sample demonstrates the ability to translate technically complex Web3 infrastructure into clear website copy while maintaining factual discipline and avoiding unsupported claims.